

Comenius University in Bratislava
Faculty of Mathematics, Physics and Informatics

Machine Learning for Generation of Graph of Given Degree and Girth

Bachelor Thesis

2026

Vladimír Jančár

Comenius University in Bratislava
Faculty of Mathematics, Physics and Informatics

Machine Learning for Generation of Graph of Given Degree and Girth

Bachelor Thesis

Study Programme: Applied Computer Science
Field of Study: Computer Science
Department: Department of Applied Informatics
Supervisor: Mgr. Ján Pastorek

Bratislava, 2026
Vladimír Jančár



THESIS ASSIGNMENT

Name and Surname: Vladimír Jančár
Study programme: Applied Computer Science (Single degree study, bachelor I. deg., full time form)
Field of Study: Computer Science
Type of Thesis: Bachelor's thesis
Language of Thesis: English
Secondary language: Slovak

Title: Machine learning for generation of graph of given degree and girth.

Annotation: Generating graphs with prescribed structural properties is a key task in combinatorial optimization, network design, and bioinformatics. Traditional algorithmic approaches (e.g., configuration model) often fail to find graphs with high girth and specific degree because the space of possible solutions is vast and the constraints are nontrivial.

Aim: The goal is to investigate potential applications of machine learning to problems in algebraic graph theory. The student will have the opportunity to engage with the state-of-art research in machine learning and algebraic graph theory and contribute to the field by generating new record graphs and training neural networks to predict algebraic properties.

Literature: Morris, C., Ritzert, M., Fey, M., Hamilton, W. L., Lenssen, J. E., Rattan, G., & Grohe, M. (2019). Weisfeiler and Leman Go Neural: Higher-Order Graph Neural Networks. *Proceedings of the AAAI Conference on Artificial Intelligence*, 33(01), 4602–4609. [<https://doi.org/10/ggfn97>] (<https://doi.org/10/ggfn97>)
Paolino, R., Maskey, S., Welke, P., & Kutyniok, G. (2024). *Weisfeiler and Leman Go Loopy: A New Hierarchy for Graph Representational Learning*.
Wu, L., Cui, P., Pei, J., & Zhao, L. (Eds.). (2022). *Graph Neural Networks: Foundations, Frontiers, and Applications*. Springer Nature Singapore. [<https://doi.org/10.1007/978-981-16-6054-2>] (<https://doi.org/10.1007/978-981-16-6054-2>)

Supervisor: Mgr. Ján Pastorek
Department: FMFI.KAI - Department of Applied Informatics
Head of department: doc. RNDr. Tatiana Jajcayová, PhD.
Assigned: 25.02.2025

Approved:

Guarantor of Study Programme

Student

Supervisor

Acknowledgments:

I would like to thank my supervisor, Mgr. Ján Pastorek, for our weekly meetings and for the discussions of my progress throughout the work on this thesis. I also thank Samuel Varchol, who works under the same supervision on a different task, for the ideas we shared.

This work was supported by the use of computational resources of the supercomputer PERUN, operated by the Supercomputing Centre at the Technical University of Košice (TUKE), Slovakia with the support of the European Union from the funds of the Recovery and Resilience Plan of the Slovak Republic within the framework of project No. 17I03-04-P03-00001, Development and design of a supercomputer for the National Supercomputing Center.

I used the AI assistant Claude (Anthropic) throughout the work on this thesis. All research ideas and designs are my own. I used it mainly to prototype them in code, so that unpromising approaches could be discarded quickly, and to take on peripheral engineering tasks such as running and monitoring computations, version control, typesetting the text and pseudocode in Typst, and designing the accompanying website. It even helped me work out how and where to write this very acknowledgment. I did not use it to generate any results, proofs, or conclusions, and I take full responsibility for the final text, in accordance with Rector's Directive No. 2/2024 of Comenius University.

The source code for this thesis is available at github.com/Vl4dk0/gnn, and an accompanying interactive website at vladimirjancar.sk.

Abstract

[One- to two-sentence framing of the overall goal: machine-learning methods for generating graphs of prescribed degree and girth, i.e. (k, g) -graphs and cage candidates.]

[Paragraph on the preparatory subtasks: which graph-property prediction tasks were studied (degree prediction, minimum-cycle / girth-related prediction), which GNN architectures were compared, and what the headline finding of these preparatory experiments is.]

[Paragraph on the main construction approaches that were tried: direct reinforcement-learning graph editing, voltage graph lifts with non-learned search, and voltage graph lifts with learned policies. Briefly state what each approach contributes.]

[Closing sentence stating the main experimental conclusion and what it implies about the role of GNNs in this setting. To be finalized once the final wording of the main claim is fixed.]

Keywords: [keywords to be finalized, current working set: graph neural networks, algebraic graph theory, cage problem, (k, g) -graphs, voltage graph lifts, heuristic search]

Contents

Abstract	ii
1 Introduction	1
2 Background and Definitions	2
2.1 Graphs	2
2.2 (k, g) -Graphs and Cages	2
2.3 Groups	3
2.4 Tabu Search	3
2.5 Graph Neural Networks	3
3 Related Work	5
3.1 Classical and Computational Cage Construction	5
3.2 Graph Neural Networks and Their Limits	6
4 Probing GNNs on Graph Properties	8
4.1 Data Generation	8
4.2 Implementation	8
5 Construction Methods for (k, g) -Graphs	10
5.1 From Guided Search to Reinforcement Learning	10
5.2 Direct Reinforcement Learning	11
5.3 Curriculum and Reward Design	11
5.4 Limits of the Direct Formulation	12
5.5 Implementation	12
6 Voltage Graph Lifts	15
6.1 Parameters of the Construction	15
6.2 Building the Lift	15
6.3 Computing Girth on the Base Graph	16
6.4 Search Methods over Voltage Assignments	19
6.5 Implementation	19
7 Refinement and Excision	22
7.1 Edge-Swap Refinement	22
7.2 Excision	25
7.3 Implementation	25
8 Forge	28
8.1 Implementation	28
9 Results	29
9.1 Predicting Degree and Minimum Cycle	29

9.2 Constructing (k, g) -Graphs	31
9.3 Summary	34
10 Future Work	36
11 Conclusion	37
Bibliography	38

List of Figures

Figure 1	A voltage lift of K_4 over Z_3	16
Figure 2	How base cycles determine the girth of the lift.	17
Figure 3	A 2-swap.	22
Figure 4	A 2-swap removing a short cycle.	23
Figure 5	A 3-swap.	23
Figure 6	Per-node accuracy by architecture: degree vs. minimum cycle.	31
Figure 7	Cage construction: coverage versus graph size relative to the Moore bound. . . .	33
Figure 8	Forge by producer: success rate and mean graph size vs. the Moore bound. . . .	34

List of Tables

Table 1	Degree prediction: best model per architecture, sorted by accuracy.	30
Table 2	Minimum-cycle prediction: best model per architecture, sorted by accuracy. . . .	30
Table 3	Cage construction by method (22 targets, 100 seeds).	32
Table 4	Forge by voltage producer, full pipeline, over all targets and seeds.	34

1 Introduction

Graph neural networks have become the standard tool for machine learning on graphs, with successes ranging from molecular property prediction to the structure models behind AlphaFold. This thesis investigates whether they can also contribute to algebraic graph theory, and in particular to the construction of record graphs of prescribed degree and girth. The objects of interest are (k, g) -graphs: graphs in which every vertex has degree k and every cycle has length at least g . A smallest such graph is a *cage*, and determining its order is a classical extremal problem [1]. The definition fits in one line, yet building small (k, g) -graphs is very hard, and the best known orders are still being improved by a mixture of hand constructions and large computer searches [2].

Bringing graph neural networks to this problem is not as straightforward as it first appears. A standard message-passing network learns from local neighborhoods, and there is a precise sense in which it cannot perceive girth at all: its expressive power is bounded by the one-dimensional Weisfeiler–Leman test [3], and girth is provably beyond what such a model can compute [4]. A network that cannot even see the property that defines the problem cannot be expected to produce or certify a high-girth graph on its own. The realistic question is therefore narrower: not whether a network can produce a cage by itself, but whether learning can *steer* where the search looks while exact algorithms decide which graphs are valid.

The recurring lesson from successful cage constructions is to constrain the search space mathematically and then explore only the freedom that remains. This thesis follows the same principle. It first probes what graph neural networks can and cannot learn about the two constraints that define a (k, g) -graph, and then works through a sequence of construction formulations of increasing structure, from unrestricted edge editing to algebraic voltage lifts in which regularity is built in and only girth is left to control. At each step the same question is asked: whether the learned component actually improves the search over its non-learned counterpart. The aim is to locate where, if anywhere, learning earns its place beside structured search rather than to claim that it replaces it.

2 Background and Definitions

2.1 Graphs

A *graph* is a pair $G = (V, E)$, where V is a finite set of vertices and E is a set of edges. In the main construction setting, graphs are simple and undirected.

The *degree* of a vertex $v \in V$, denoted $\deg(v)$, is the number of edges incident with v . The *open neighborhood* of v , denoted $N(v)$, is the set of vertices adjacent to v .

A graph is k -regular if every vertex has degree k .

A *walk* is a sequence of vertices in which consecutive vertices are adjacent.

A *path* is a walk with no repeated vertices.

A *cycle* is a closed path of length at least 3.

A *closed walk* is a walk that starts and ends on the same vertex.

A closed walk is *reduced* if it never traverses an edge immediately followed by the same edge in reverse.

The *girth* of a graph, denoted $\text{girth}(G)$, is the length of its shortest cycle. If a graph has no cycles, its girth is treated as infinite.

A vertex v is *active* if $\deg(v) > 0$. The *active subgraph* of G is the subgraph induced by the active vertices, i.e. G with all isolated vertices removed. This is convenient when a construction algorithm carries a large fixed vertex pool but uses only the vertices it has already drawn into edges.

2.2 (k, g) -Graphs and Cages

A (k, g) -graph is a k -regular graph whose girth is at least g . A (k, g) -cage is a smallest k -regular graph of girth g . Its order is commonly denoted $n(k, g)$. In search algorithms it is convenient to test the condition $\text{girth}(G) \geq g$, because this accepts every valid candidate for a target lower bound on girth.

The Moore bound gives a general lower bound on the order of a (k, g) -graph. It is useful both as a theoretical benchmark and as a practical target size for construction algorithms.

For $k \geq 2$ and $g \geq 3$, the Moore bound is defined as follows:

$$2 \mid g \rightarrow M(k, g) = 1 + k \sum_{i=0}^{\frac{g-3}{2}} (k-1)^i.$$

$$2 \nmid g \rightarrow M(k, g) = 2 \sum_{i=0}^{\frac{g-2}{2}} (k-1)^i.$$

A graph meeting this bound is called a Moore graph. Moore graphs are rare. For most parameter pairs, the practical problem is therefore to close the gap between the Moore lower bound and the best published upper bounds.

2.3 Groups

A *group* $(\Gamma, \cdot)$ is a set Γ with a binary operation $\cdot$ that combines any two elements $a, b \in \Gamma$ into an element $a \cdot b \in \Gamma$, subject to three axioms: the operation is *associative*, meaning $(a \cdot b) \cdot c = a \cdot (b \cdot c)$; there is an *identity* element $e \in \Gamma$ with $e \cdot a = a \cdot e = a$ for every a ; and every $a \in \Gamma$ has an *inverse* $a^{-1} \in \Gamma$ with $a \cdot a^{-1} = a^{-1} \cdot a = e$.

A group is *abelian* if its operation commutes, i.e. $a \cdot b = b \cdot a$ for all $a, b \in \Gamma$. Otherwise it is *non-abelian*.

The *order* of a group is the number of its elements, and a group with finitely many elements is *finite*. The *order* of an element a is the smallest positive integer m for which $a^m = e$, the m -fold product of a with itself. The constructions in this thesis use finite groups, most often the *cyclic group* Z_n : the integers $\{0, 1, \dots, n-1\}$ under addition modulo n , with identity 0 and the inverse of i equal to $n-i$ modulo n .

2.4 Tabu Search

Tabu search is a local-search metaheuristic for minimizing a cost over a space of candidate solutions. From the current solution it considers the neighbors reachable by a small move and steps to one of them, repeating until the cost can no longer be lowered. To avoid becoming trapped, it may step to a neighbor that is temporarily worse when no improving move is available.

A *tabu list* is the short-term memory that makes this possible. It records the most recently applied moves, or their reversals, and forbids them for a few steps. This prevents the search from immediately undoing its last move and oscillating between the same few solutions, and it pushes the search to explore away from a local minimum instead of sliding back into it.

2.5 Graph Neural Networks

Ordinary feed-forward neural networks expect fixed-size vectors. Graphs do not have a fixed number of vertices, and their vertices do not have a canonical ordering. Graph neural networks address this by applying the same local update rule at every vertex. This makes the model independent of the input graph size and equivariant to relabeling of vertices.

Message-passing graph neural networks process graph-structured data by iteratively updating vertex representations through layers [5]. At layer t , each vertex v has a hidden representation h_v^t . The layer aggregates messages from the neighborhood $N(v)$ and then updates the representation:

$$m_v^t = \text{AGG}(\{h_u^{t-1} : u \in N(v)\})$$

$$h_v^t = \text{UPD}(h_v^{t-1}, m_v^t).$$

Several parameters determine what such a model can express. The *number of layers* controls how far information can travel: after one layer, a vertex representation depends on its immediate neighbors, after two layers it can depend on vertices at distance two, and so on.

The *hidden dimension* controls how much information can be stored in each vertex representation. The *aggregation rule* determines which neighborhood statistics are easy to preserve, for example, normalized aggregation can make raw counting harder, while sum-like aggregation keeps count information more directly.

Edge features, positional encodings, and graph-level pooling become important when the task depends on edge labels or on the position of a vertex inside the whole graph. This is why local quantities such as degree are easier for standard GNNs than global quantities such as girth.

Grohe’s survey on the logic of graph neural networks provides the theoretical framing for this limitation [3]. Standard message-passing GNNs are closely related to color refinement and the one-dimensional Weisfeiler–Leman algorithm. This connection is important for the thesis because it explains why some graph properties are naturally accessible to ordinary GNNs while other global or symmetry-sensitive properties may require stronger architectures, additional features, or a different formulation of the search problem.

The experiments use several GNN architectures:

- *GCN*, a stable baseline based on degree-normalized neighbor aggregation [6].
- *GraphSAGE*, an inductive architecture based on learned aggregation [7].
- *GIN*, a sum-aggregation architecture with expressivity related to the Weisfeiler–Leman test [8].
- *GINE*, an edge-aware variant of GIN used when edge attributes such as voltage labels are part of the input [9].
- *GPS*, a graph transformer architecture combining local message passing with global attention [10]. GPS models often benefit from structural or positional encodings, such as random-walk-based encodings, because attention by itself does not identify a vertex’s structural role in the graph.
- *Loopy GNN*, a cycle-aware architecture relevant to minimum-cycle and girth-related tasks [11].

3 Related Work

The work builds on two lines of research: graph-theoretic construction of small regular graphs, and learned heuristics for difficult discrete search problems. The relevant background is not every known cage construction, but the recurring pattern behind successful methods: restrict the search space mathematically, then use computation to explore the remaining choices.

3.1 Classical and Computational Cage Construction

The cage problem has been studied through algebraic constructions, finite geometries, Cayley graphs, voltage graph lifts, exhaustive generation, local search, and excision. The Dynamic Cage Survey provides the broad background and record-holder context [1].

The most relevant lesson from classical construction is that successful methods rarely search over arbitrary graphs. They usually impose structure first and search only inside the remaining degrees of freedom. Algebraic families illustrate this directly. Incidence graphs of finite projective planes and generalized quadrangles yield $(q + 1)$ -regular graphs of girth 6 and 8 for prime powers q . Generalized hexagons give analogous constructions for girth 12. Cayley graphs on suitable groups produce highly symmetric candidates, and Ramanujan graphs [12] provide explicit algebraic families with optimal spectral properties and large girth.

These methods do not search adjacency at all. They encode the constraint into the algebraic object.

Voltage graph lifts continue the same pattern. A small base graph and a finite group define a covering graph whose regularity is inherited from the base. Exoo and Jajcay formalize the relation between voltages on closed walks in the base graph and the girth of the lift [13]. This reduces the search from $O(n^2)$ adjacency choices on a lifted graph to a much smaller voltage assignment on the base.

A second family of methods relies on local search inside the space of k -regular graphs of a fixed order. Hill climbing and tabu search use edge swaps that preserve regularity and minimize a cycle-based cost function. Tabu lists prevent the search from immediately reversing recent moves. Excision methods take a known small cage, remove a carefully chosen subgraph (typically a tree or a small structured neighborhood), and reconnect the remaining deficient vertices to obtain a regular graph one to several vertices below the previous best [14], [15]. Constraint programming and integer programming encodings have also been studied, but the size of the target graphs limits their direct applicability.

Recent computational work combines several of these ideas. Exoo et al. couple voltage-lift generation with advanced tabu search, hill climbing using a distance distribution cost, and tree excision, and report eleven new upper bounds on cage orders in a single pipeline [2]. The takeaway for this thesis is methodological: improvements come from chaining structurally constrained search methods, not from any single algorithm.

These methods provide the natural non-learning baselines. A learned method is not very informative if it only beats unrestricted random edge editing. A stronger comparison asks whether learning improves a search process that already uses graph-theoretic structure, such as a voltage-lift search with exact short-cycle checks.

3.2 Graph Neural Networks and Their Limits

The machine-learning side is closest to work on graph neural networks and learned heuristics for combinatorial search, not only to graph generation. Work such as NeuroSAT showed that neural networks can learn useful signals on symbolic search instances even when they do not replace the exact solver [16]. The analogy here is that a learned component may guide which states or moves are explored first, while exact algorithms still verify the final graph.

Standard message-passing GNNs are well-suited to permutation-equivariant scoring of graph states [3]. Garg, Jegelka, and Jaakkola prove that local message-passing GNNs cannot compute several global graph invariants, including girth and diameter, and give the first data-dependent generalization bounds for such models [4]. This is precisely the regime of the cage problem: high girth, vertex-transitive or nearly so, and locally tree-like. On such graphs, L -layer message-passing GNNs collapse to nearly identical vertex embeddings, and standard architectures cannot resolve the global structure required to verify or score a candidate construction.

Two responses to this limit are relevant here. Subgraph-based architectures add explicit cycle or motif counts as structural features, which lifts expressivity above plain 1-WL message passing [17]. Cycle-aware hierarchies such as the Loopy framework [11] extend message passing with cycle bases. Both lines of work are used in this thesis as candidate architectures when the prediction target depends on cycle information.

Degree regularity, on the other hand, can be enforced by construction, and short-cycle violations can often be detected exactly. This motivates learned scoring and move ranking rather than replacing the mathematical verification step.

Reinforcement learning has also been applied directly to extremal graph problems. Wagner combined a cross-entropy method with a small neural policy to construct counterexamples to combinatorial conjectures, on graphs of at most a few tens of vertices with a single

scalar objective [18]. The later RLGT framework formalizes the same setting with structural rewards [19].

Where reinforcement learning has succeeded on related problems, the action space already carried structure: Freire et al. apply it to the LDPC components of hypergraph product codes, where the search operates on a constrained construction rather than on arbitrary adjacency [20]. These results frame the construction experiments in this thesis — the unrestricted edge-editing formulation is the hardest case, and the algebraic formulations deliberately reintroduce structure.

4 Probing GNNs on Graph Properties

Degree prediction and minimum-cycle prediction serve as controlled probes for the two structural constraints that define a (k, g) -graph. In both tasks the graph is given and the model assigns a label to each vertex.

Degree prediction is local and exactly verifiable. A vertex degree can be recovered from the immediate neighborhood, so this task tests whether the model, data pipeline, and training loop can learn a simple structural quantity. If an architecture fails to predict degree reliably, it is not a good candidate for harder construction tasks without further tuning.

Minimum-cycle prediction is a harder structural task. For each vertex, the target is the length of the shortest cycle containing that vertex, with target 0 for vertices that do not lie on any cycle. Unlike degree, this cannot be determined from immediate neighbors alone. It requires information about paths that leave a vertex and later return to it. This makes the task closer to the girth constraint used in constructing (k, g) -graphs.

How the standard graph neural network architectures compare on these two probes is reported in Chapter 9.

4.1 Data Generation

Both tasks use dynamically generated random graphs. Instead of training on a fixed set of graphs, new Erdos–Renyi graphs are sampled during training. This reduces the chance that a model simply memorizes a finite training set and gives a cheaper way to expose it to many small graph structures.

Labels are generated by exact graph algorithms. Degree labels are computed as $\deg(v)$ for each vertex. Minimum-cycle labels are computed by searching for the smallest cycle containing each vertex. This keeps the supervision reliable even though the graphs themselves are random.

The input features are intentionally simple. The experiments use either a constant one-dimensional feature or a four-dimensional feature vector consisting of a normalized node index, two random real-valued coordinates, and one additional random scalar. These features should not be interpreted as meaningful graph attributes. They mainly give the neural network a vertex-level input tensor while forcing the model to learn from graph structure.

4.2 Implementation

Training samples a fresh random graph at every step and labels it exactly, rather than reusing a fixed dataset. Degree labels are immediate, but the minimum-cycle label is the interesting one.

For a vertex v we look, through each incident edge in turn, for the shortest way back to v that avoids that edge. Removing the edge to a neighbor u and finding a shortest path from u back to v gives a cycle through v one longer than that path; the shortest over all neighbors is the label, and a vertex on no cycle is labeled 0.

```
func min_cycle_label(G: Graph, v: Vertex):
    best: Integer = infinity;
    for each neighbour u of v in G:
        remove edge (v, u) from G;
        // the shortest detour back to v closes a cycle through this
edge
        d: Integer = length of the shortest path from u to v in G;
        if d is finite:
            best = min(best, d + 1);
        restore edge (v, u) to G;
        // a vertex on no cycle has no such detour
    if best is infinity:
        return 0;
    return best;
```

Algorithm 1: Exact minimum-cycle label for a single vertex.

5 Construction Methods for (k, g) -Graphs

The target problem is to construct small (k, g) -graphs. This is harder than predicting a graph invariant, because the algorithm must actively choose edges while keeping the partial object close to a valid regular graph. The search space grows quickly with the number of vertices, and most arbitrary edge choices lead to invalid or unpromising graphs.

The problem is therefore more naturally viewed as a search problem than as a single prediction problem. A construction algorithm repeatedly holds a partial object, chooses a modification, checks which constraints are still satisfiable, and continues until it either reaches a valid graph or enters an unpromising region of the search space. Machine learning can enter this process in several different ways: as a learned heuristic for deterministic search, as a policy trained from complete construction attempts, or as a scoring function inside a constrained algebraic search.

5.1 From Guided Search to Reinforcement Learning

A natural first idea is to use deterministic search and let a GNN guide it. For example, one can build a graph edge by edge and use a learned scoring function to prioritize partial graphs that appear more likely to extend to a valid (k, g) -graph.

The difficulty is not the search loop itself. The difficulty is that ordinary supervised training does not apply: there is no straightforward way to label partial states well enough to train a model to predict whether to expand them. Positive examples can be obtained by taking subgraphs of known valid graphs. Useful negative examples are harder: a partial graph may look bad but still be extendable through further edits. Deciding whether a partial graph can be completed into a valid (k, g) -graph leads back to the same kind of problem. A supervised heuristic can therefore end up training mostly on easy examples far from the decision boundary.

This explains why a purely supervised A-star-style approach was not pursued as the main construction method.

The attractive scenario would be a good heuristic that guides A-star-style search without a pre-built dataset. Reinforcement learning provides exactly this: instead of asking for a dataset that says whether each partial graph is extendable, the model learns from complete construction attempts. The environment supplies rewards based on the observed consequences of actions: whether an edge edit was valid, whether the graph moved closer to regularity, whether short cycles were avoided, and whether the final graph satisfies the target constraints.

In this sense, the reinforcement-learning approach is not separate from search. It is a learned version of heuristic search. A hand-designed search algorithm chooses moves

according to a manually specified rule. An RL policy attempts to learn such a rule from repeated interaction with the construction environment.

It is important to keep in mind what this formulation does and does not provide. Individual illegal moves can be ruled out, as the next section describes: an addition that exceeds degree k or closes a cycle shorter than g is simply never offered to the agent, so every intermediate graph stays legal. What no such rule provides is completion. Unlike the algebraic constructions of the next chapter, where regularity is automatic, nothing here guarantees that a sequence of legal edits ever reaches a graph that is k -regular on enough vertices. The agent must learn to assemble one edge by edge rather than stall in a legal but unfinishable state.

The direct-RL formulation is therefore useful partly as a diagnostic: it tests how far a learned search can go when local legality is enforced but global completion is left to the policy.

5.2 Direct Reinforcement Learning

In the direct reinforcement-learning formulation, construction is represented as a sequential decision problem and trained with Proximal Policy Optimization [21]. The state is a partially constructed graph. Each action selects an unordered pair of vertices. If the edge is absent, the action attempts to add it. If the edge is present, the action attempts to remove it.

This direct graph-editing formulation has an advantage: it does not require a precomputed dataset of labeled partial graphs. The model learns from interaction. However, it also makes the action space and the reward design essential, and these are addressed below.

Several invalid or unhelpful actions are removed from the action space before the policy sees them, so the agent cannot select them at all:

- An edge cannot be added if either endpoint already has degree k .
- An edge cannot be added if it would create a cycle shorter than g .
- An edge cannot be removed if doing so would disconnect the active subgraph.
- An edge cannot be removed if it would leave too few active vertices to reach the Moore-bound target.

5.3 Curriculum and Reward Design

The motivation for the curriculum is direct: at uniform difficulty an untrained agent almost never produces a valid graph, so the learning signal collapses and nothing is learned. Starting from parameter pairs with a small Moore bound raises the per-episode success rate enough that the agent has something to learn from, and the difficulty can then be increased once that signal is reliable.

The environment orders parameter pairs by their Moore bound and unlocks the next pair only after the agent solves at least half of its last eight episodes on the current pair. Training starts from $(3, 5)$.

Curriculum learning is what makes the direct RL formulation learn at all. Without it, the agent never produces a valid $(3, 5)$ -graph. With it, the agent reaches and solves the first stages, the $(3, 5)$ -graph and the $(3, 6)$ -graph. Reaching even these stages also depended on the reward design described next.

The environment uses a shaped reward. Earlier reward designs (in my codebase, not in related work) concentrated almost all signal at the end of an episode, and successful episodes were too rare for that signal to guide learning reliably. A per-step progress term was added so that the agent receives a positive signal when an action moves the partial graph closer to a valid construction and a negative one when it moves away. A valid final graph receives a large terminal reward, but the per-step term carries most of the learning signal.

The progress term is implemented as a potential-based shaping signal [22]. Let s denote a state (a partial graph) and let $\Phi(s)$ denote its *potential*: a scalar that grows as the active subgraph approaches a k -regular graph on the Moore-bound number of vertices, rewarding both each active vertex approaching degree k and the total edge count approaching the corresponding $kM/2$ edges.

The per-step shaping reward is

$$\Phi(s') - \Phi(s),$$

where s' is the state reached after taking the action in state s . The value is positive when the partial graph moves closer to the target and negative when it moves away.

5.4 Limits of the Direct Formulation

The direct graph-editing formulation is useful as a baseline, but it has structural weaknesses. Regularity is not guaranteed. It must be learned or enforced through action pruning and rewards. The action space also grows quadratically with the number of available vertices: on a graph with n vertices, there are $\binom{n}{2}$ possible unordered pairs, and each pair can become an add-or-remove decision depending on the graph state.

These limits motivate the voltage-lift formulation of the next chapter, which builds regularity into the representation rather than leaving it to the agent.

5.5 Implementation

Two parts of the direct formulation need a concrete procedure: which edits the policy may make, and how each step is rewarded. The policy samples only from a precomputed set of legal edits, so an illegal one can never be selected. A missing edge is legal to add unless an

endpoint already has degree k or its endpoints are nearer than $g - 1$ apart, which would close a cycle shorter than g . A present edge is legal to remove unless that disconnects the active subgraph or drops it below the Moore-bound vertex count.

```
func legal_edits(G: Graph, k: Integer, g: Integer):
    legal: Set of Vertex Pairs = empty;
    for each unordered pair (u, v) of vertices:
        if G has edge (u, v):
            // a removal must keep the active subgraph connected and
            // keep at least Moore-bound-many active vertices
            if removing (u, v) keeps the active subgraph connected
                and leaves at least Moore-bound-many active vertices:
                legal <- (u, v);
        else:
            // an addition must not exceed degree k and must not
            // close a cycle shorter than g
            if degree(u) < k and degree(v) < k
                and distance from u to v in G is at least g - 1:
                legal <- (u, v);
    return legal;
```

Algorithm 2: Building the set of legal edits offered to the policy.

The reward for a step is the change in a potential Φ across the edit, $\Phi(s') - \Phi(s)$, where Φ rises as the active subgraph approaches a k -regular graph on the Moore-bound number of vertices. It combines two normalized scores: how far the active vertices have progressed toward degree k , and how close the edge count is to that of the finished graph. A completed valid graph adds a large terminal bonus.

```

// k is the target degree, M the Moore-bound vertex count, g the target
girth
func potential(s: State):
    G = graph of state s;
    active = vertices of G with degree > 0;
    progress = sum over v in active of min(degree(v), k);
    // degree progress toward k-regular, in [0, 1]
    reg = progress / (k * M);
    // edge count toward the finished graph, in [0, 1]
    dens = min(edges of G / (k * M / 2), 1);
    return weighted sum of reg and dens;

func reward(s: State, s': State):
    r = potential(s') - potential(s);
    // bonus only when the edit completes a valid graph
    if s' is a valid (k, g)-graph:
        r = r + large terminal bonus;
    return r;

```

Algorithm 3: The potential and the per-step shaped reward.

6 Voltage Graph Lifts

Voltage graph lifts are an alternative construction formulation in which k -regularity becomes structural: every lift of a k -regular base graph is itself k -regular. The control problem then reduces to choosing group labels on the edges of a small base graph instead of choosing every edge of the large lifted graph, and the learned components introduced later in this chapter operate inside this reduced space.

6.1 Parameters of the Construction

A voltage lift is determined by three choices, and only three. Everything else is forced once they are fixed.

The first choice is a finite group Γ . Its role is to set how many copies of the base graph the lift contains: the lift consists of $|\Gamma|$ copies of B , so a larger group produces a larger graph.

The second choice is an orientation of each base edge. This is only a bookkeeping convention for the direction in which a voltage is read. Reversing an arc and replacing its voltage a by the inverse a^{-1} produces exactly the same lift.

The third choice is the voltage assignment α , which labels each oriented edge with an element of Γ . This is the only choice that affects girth. It is precisely the voltage assignment that the search and learning procedures set out to optimize.

What is not free follows from these three. The number of vertices is $|V(B)| \cdot |\Gamma|$, the degree of every lift vertex equals the degree of the base vertex it copies, and the adjacencies of the lift are fixed by the rule in the next subsection.

6.2 Building the Lift

The base graph B is small and may contain loops and parallel edges. For undirected graphs each edge is represented by two opposite arcs. If one direction carries voltage a , the reverse carries a^{-1} .

The lift has vertex set $V(B) \times \Gamma$. Writing (v, h) for the copy of base vertex v on layer $h \in \Gamma$, an arc from u to v with voltage a contributes, for every group element h , the edge

$$(u, h) \text{ to } (v, ha).$$

This thesis uses right multiplication by the voltage element. For abelian groups the order of multiplication is invisible, but it matters for non-abelian groups.

As an example, take $B = K_4$ and $\Gamma = Z_3$, so the lift has $4 \cdot 3 = 12$ vertices. Figure 1 shows an assignment in which a spanning star carries voltage 0 and the remaining triangle carries voltage 1. An edge with voltage 1 from u to v produces $(u, h) - (v, h + 1)$ for $h \in \{0, 1, 2\}$, so each base edge becomes three lift edges and the three layers are stitched together cyclically. The resulting graph is 3-regular and connected.




Figure 1: A voltage lift of K_4 over $\mathbb{Z}_3$.

Regularity is structural. The reason the construction guarantees k -regularity is direct. Fix a base vertex u of degree k and a layer h . Each arc leaving u contributes exactly one lift edge at (u, h) : the arc to base neighbor v with voltage a produces the single edge $(u, h) - (v, ha)$. The k arcs at u therefore yield exactly k edges at (u, h) , regardless of which voltages were chosen, so every lift vertex has degree k .

This is the property that makes the formulation attractive for (k, g) -graphs. In direct edge-by-edge generation the agent must continually maintain the degree of every vertex. Here the degree constraint holds automatically, and only the girth constraint remains to be controlled through the voltage assignment.

6.3 Computing Girth on the Base Graph

The second advantage is that girth can be read off the base graph without building the lift. Consider a reduced closed walk W in the base graph. Multiplying the voltages along its arcs (using a^{-1} for an arc traversed backward) gives the *net voltage* of W .

Following the corresponding lift edges around W from layer h returns to the starting base vertex, but on layer hs , where s is the net voltage. Two cases arise:

- If s is the identity, the walk closes after one trip and lifts to a cycle of the same length as W .
- If s is not the identity, the walk must be repeated until the accumulated layer shift returns to the identity. The number of repetitions is the *order* of s , so W lifts to a cycle of length $|W| \cdot \text{ord}(s)$.

The girth of the lift is the length of the shortest cycle obtained this way:

$$\text{girth}(\text{lift}) = \min_W |W| \cdot \text{ord}(s),$$

minimized over reduced closed walks W of the base graph, where s is the net voltage of W — equivalently, the shortest base walk whose net voltage is the identity.

Figure 2 illustrates both cases on the running example. Every triangle of K_4 has net voltage 1, which has order 3 in Z_3 . A triangle therefore cannot close after one trip and instead wraps three times into a single 9-cycle that visits all three layers.

The 4-cycle 0–2–3–1, by contrast, has net voltage 0 and survives as an ordinary 4-cycle. It is the shortest surviving cycle, so the lift has girth 4, raising the girth of 3 in K_4 . No voltage assignment on K_4 over Z_3 reaches girth 5. Higher girth requires a larger group or a different base graph, which is why the search later ranges over many base–group pairs.




Figure 2: How base cycles determine the girth of the lift.

This analysis gives a cheap exact cost for search: enumerate the reduced closed walks shorter than the target girth g and count those with identity net voltage. The count is zero exactly when the lift has girth at least g , and it is computed on the small base graph rather than on the full lift. This exact cost is used by the search and reinforcement-learning procedures below, and it is the main reason the learned component can sit inside a constrained search instead of generating an entire graph from scratch.

Reducing the search space. The search becomes cheaper when fewer edges need a voltage: exhaustive or guided search over m undirected base edges scales with $|\Gamma|^m$, so any voltage that can be fixed in advance, without ruling out a single lift, shrinks the problem. Several base voltages can indeed be fixed for free, and this subsection explains which ones and why.

The reason such fixing is harmless is that the labels attached to the copies of a base vertex are arbitrary. At a vertex v , the $|\Gamma|$ copies may be renamed among themselves. This is a pure relabeling of lift vertices, so the lifted graph is unchanged up to isomorphism.

Renaming the copies of v by a group element $\tau(v)$ shifts the voltage of every arc touching v : it adds τ at the arc's tail and subtracts it at the arc's head, so an arc from u to v acquires the new voltage

$$\alpha(u \rightarrow v) + \tau(u) - \tau(v).$$

Choosing one such element $\tau(v)$ at every vertex, and letting it act on the voltages in this way, is called a *gauge transformation*.

A gauge transformation cannot change the girth, because it leaves the net voltage of every closed walk untouched. Adding up the corrections along a closed walk, each vertex it passes through is entered along one arc — contributing $-\tau$ — and left along the next — contributing $+\tau$. The two cancel. Over the whole walk the corrections sum to $\tau(\text{start}) - \tau(\text{end}) = 0$. Since girth is determined entirely by the net voltages of closed walks, the lift is the same graph before and after.

This freedom can now be spent to drive voltages to zero. Forcing a single arc to zero means choosing the potentials so that

$$\alpha(u \rightarrow v) + \tau(u) - \tau(v) = 0.$$

The catch is that not every edge can be zeroed at once, and the reason is exactly the cancellation above: around any cycle the τ corrections sum to zero, so the net voltage of a cycle is fixed — a gauge transformation cannot touch it. If a cycle carries a nonzero net voltage, its edges therefore cannot all be made zero. Cycles are the obstruction, and only a set of edges that contains no cycle can be zeroed simultaneously.

This is what singles out a spanning tree. A spanning tree is the largest set of edges of a connected graph that contains no cycle, and it has $|V(B)| - 1$ edges. Its voltages can always be cleared: root the tree, set $\tau = 0$ at the root, and work outward.

Each tree edge joins an already-fixed vertex u to a new vertex v , and the choice $\tau(v) = \alpha(u \rightarrow v) + \tau(u)$ makes that edge's voltage zero. Because the tree has no cycle, every vertex is reached exactly once, along a unique path from the root, so these choices never conflict. Adding any further edge would close a cycle, whose net voltage is beyond reach, so a spanning tree is as far as the zeroing can go.

What remains are the $m - |V(B)| + 1$ non-tree edges — the only edges that carry genuine freedom. The effect is substantial. For the example, K_4 has six edges and a spanning tree of

three, so three voltages are fixed to zero and only three remain to be chosen: the triangle on vertices 1, 2, 3 in the assignment above. The search assigns three voltages instead of six.

6.4 Search Methods over Voltage Assignments

Several search strategies share the voltage-lift formulation. They differ in how they choose voltage assignments, but they all keep exact verification at the end: a generated graph is accepted only after connectedness, regularity, and girth are checked.

For small groups and few base edges, exhaustive search enumerates every voltage assignment. This is only practical when $|\Gamma|^m$ is small, where m is the number of undirected base edges. Random search samples assignments and verifies them. It is simple, but it is an important baseline because some small cases are easy enough that random search succeeds.

The tabu search follows recent computational work on small regular graphs [2]. It fixes tree-edge voltages to the identity, so only non-tree edges are searched. It then changes one voltage at a time and minimizes the number of short identity-voltage walks. A tabu list prevents the algorithm from immediately undoing recent moves.

The GNN-guided beam search assigns voltages edge by edge. A trained girth-predictor network scores partial assignments, and the search keeps only the highest-scoring candidates. *[Architecture and evaluation of this predictor to be detailed in Chapter 9.]* A meta-search procedure additionally enumerates choices of base graph and finite group. For cubic graphs, the candidates include the dumbbell base, a four-node cubic base, the prism base, and a Petersen-like base. Candidate groups include cyclic groups, dihedral groups, direct products, and semidirect products.

The reinforcement-learning environment treats one voltage assignment as one episode. The state is the base graph with partial voltage labels, the action is a group element, and the episode length is the number of base edges. This is much shorter than the direct edge-editing environment. The policy model uses either a GINE-only encoder or a GPS encoder with GINE as the local message-passing component, followed by global pooling, an actor head over group elements, and a critic head for PPO.

6.5 Implementation

Because the search uses cyclic groups Z_n , a voltage is simply an integer in $\{0, \dots, n-1\}$, so a voltage assignment is a dictionary mapping each base edge to an integer, with the backward direction of an edge contributing the negated voltage modulo n . Two routines carry the chapter. The first reads the girth of a lift directly off the base graph: from every start vertex it walks reduced closed walks, adding edge voltages modulo n , and the shortest walk that returns to its start with net voltage 0 is the shortest cycle of the lift.

```

func lift_girth(base: Graph, voltage: Dictionary of Edge to Integer,
               n: Integer, limit: Integer):
    best: Integer = infinity;
    for each start vertex of base:
        // each entry: current vertex, net voltage, walk length, last
        arc
        stack: Stack of (Vertex, Integer, Integer, Arc) = empty;
        stack <- (start, 0, 0, none);
        while stack is not empty:
            node, net, length, last = top of stack;
            remove top of stack;
            // a closed walk with net voltage 0 lifts to a cycle
            if node is start and net is 0 and length >= 3:
                best = min(best, length);
                continue;
            if length >= limit:
                continue;
            for each arc a leaving node:
                // stay reduced: never retrace the previous arc
                if last is not none and a is the reverse of last:
                    continue;
                step: Integer = voltage[edge of a];
                // the backward direction carries the negated voltage
                if a runs against its edge:
                    step = (n - step) mod n;
                stack <- (head of a, (net + step) mod n, length + 1, a);
    return best;

```

Algorithm 4: Girth of the lift, computed by walk enumeration on the base graph.

The second routine is the tabu search that varies a voltage assignment toward the target girth. It fixes the spanning-tree edges to 0 and moves only the free, non-tree edges. Its cost reuses the same walk enumeration: the number of closed walks shorter than g with net voltage 0, which is exactly the number of short cycles that would appear in the lift. It changes one free voltage at a time to drive that count to 0, and a short tabu list keeps it from undoing a recent move.

```

// cost(voltage) = number of closed walks shorter than g with net
voltage 0,
func voltage_tabu(base: Graph, n: Integer, g: Integer, budget: Integer):
  free: List of Edges = the non-tree edges of base;
  // tree edges stay 0 (a fixed gauge); only free edges are searched
  voltage: Dictionary of Edge to Integer = 0 on every edge;
  for each edge e in free:
    voltage[e] = a random integer in 0 .. n - 1;

  tabu: Queue of (Edge, Integer) = empty;
  repeat up to budget:
    if cost(voltage) is 0:
      return voltage;

    // try every single-edge change and keep the one that helps most
    best_edge: Edge = none;
    best_value: Integer = 0;
    best_cost: Integer = infinity;
    for each free edge e:
      old: Integer = voltage[e];
      for each value c in 0 .. n - 1 with c != old:
        // a move is the pair (edge, new value); skip forbidden
        if (e, c) is in tabu:
          continue;
        voltage[e] = c;
        trial: Integer = cost(voltage);
        if trial < best_cost:
          best_cost = trial;
          best_edge = e;
          best_value = c;
        // undo the trial before testing the next value
        voltage[e] = old;

    // forbid undoing this move for a while, then commit it
    tabu <- (best_edge, voltage[best_edge]);
    voltage[best_edge] = best_value;
    drop the oldest tabu entry if the list is too long;
  return voltage;

```

Algorithm 5: Tabu search over free voltages, driving short cycles to zero.

7 Refinement and Excision

This chapter describes two methods that act on an already-constructed k -regular graph rather than building one from scratch. Edge-swap refinement edits a graph to remove short cycles while preserving its degree, and excision removes vertices from an existing graph to obtain a smaller one.

7.1 Edge-Swap Refinement

A voltage lift always returns a k -regular graph, but its girth may fall short of the target g , since a few short cycles can survive the construction. Such a candidate is close to a solution, and discarding it wastes the work already spent. The refinement step instead edits a small number of edges to remove the surviving short cycles while keeping the graph k -regular.

This reverses the difficulty seen during construction. There, regularity was automatic and girth was the hard constraint; here the graph already has the right degree, so the edit must improve girth without breaking regularity.

The basic move is the 2-swap. Take two edges with no shared endpoint, say A–B and C–D on four distinct vertices, delete both, and reconnect the same four vertices the other way, for example as A–D and B–C (Figure 3). Each of the four vertices loses one edge and gains one, so every degree is preserved and a k -regular graph stays k -regular. Four vertices joined by two disjoint edges can be connected in three ways, so from the original pairing there are two alternative 2-swaps available.




Figure 3: A 2-swap.

The two disjoint edges A–B and C–D (left, dashed) are replaced by A–D and B–C (right). Every vertex keeps its degree, so regularity is preserved.

A 2-swap can destroy a short cycle. To remove a triangle $A-B-C$, swapping its edge $A-B$ against a disjoint edge $D-E$ deletes $A-B$ and breaks the triangle, again without changing any degree (Figure 4).




Figure 4: A 2-swap removing a short cycle.

The triangle $A-B-C$ and a disjoint edge $D-E$ (left) are swapped to $A-D$ and $B-E$ (right): the edge $A-B$ is gone, so the triangle is broken, and all degrees are intact.

A larger move, the 3-swap, edits three edges at once. Take three pairwise-disjoint edges, on six distinct vertices, and rewire them into a different degree-preserving matching on the same six vertices (Figure 5). Like the 2-swap it leaves every degree unchanged, but it reaches configurations that no single 2-swap can produce.




Figure 5: A 3-swap.

Three disjoint edges $A-B$, $C-D$, $E-F$ (left) are rewired to $A-C$, $B-E$, $D-F$ (right). All six vertices keep their degree.

A swap can remove one short cycle while creating another elsewhere, so these moves cannot be applied blindly. They have to be driven by a search that recognizes which swaps make progress.

Search. Progress means having fewer short cycles, so the search measures exactly that. Write $N_c(G)$ for the number of cycles of length c in G . The simplest cost is the total number of cycles shorter than the target,

$$\sum_{c=3}^{g-1} N_c(G),$$

which drops by one whenever a short cycle is removed and reaches 0 exactly when no cycle shorter than g remains, the girth target. Because it counts every short cycle, the cost responds to each swap that removes one, giving the search feedback at every step.

Not all short cycles are equally harmful, however. A triangle must be removed before the girth can exceed 3, whereas a cycle just below the target is almost acceptable, so the search should attack the shortest cycles first. Weighting each length so that shorter cycles count for exponentially more makes this automatic:

$$f(G) = \sum_{c=3}^{g-1} w_c \cdot N_c(G), \quad w_c = 2^{g-c}.$$

A triangle then outweighs any single longer cycle, and the cost still reaches 0 exactly at the girth target.

The search minimizes f by repeatedly applying the swap that lowers it most, using tabu search so that it does not stall. A tabu list forbids reversing the moves just made, which stops the search oscillating between two graphs and lets it accept a temporarily worse swap to climb out of a local minimum without being pulled straight back.

The routine move is the 2-swap; when no 2-swap improves the cost and short cycles remain, the search escalates to a 3-swap, which can cross a barrier that traps the smaller move. Because there are $O(|E|^3)$ candidate 3-swaps against $O(|E|^2)$ 2-swaps, the escalation is occasional rather than routine.

Scoring candidate swaps is the cost bottleneck. There are $O(|E|^2)$ possible 2-swaps each iteration, and evaluating one exactly means recounting the short cycles it changes. A trained evaluator network replaces this with a cheaper estimate: a single pass of a graph neural network produces an embedding for every vertex, and the predicted change in cost for a swap is read from the embeddings of its few endpoints through a small head.

The expensive graph reasoning is computed once per iteration and shared across all candidate swaps, which are scored together in one batched pass. The network uses cycle-count features and random-walk positional encodings, which break the symmetry that otherwise makes vertices indistinguishable on highly regular graphs.

Because the estimate is approximate, it is used only to rank candidates: the search verifies the most promising swaps exactly before committing to one. The approximation therefore affects only how quickly and how well the search proceeds, never the validity of the result, since every accepted graph is checked exactly against the girth target. This is the same division of labor as in the construction stage, where a learned component guides the search and exact verification decides.

This refinement is the most recent component of the pipeline, and its full evaluation is still in progress. An initial training run on 100{, }000 generated swap labels reduced the mean squared error of the move evaluator by roughly 30% relative to its initialization. Held-out evaluation, a comparison against the unguided search, and feature ablations are reported in Chapter 9.

7.2 Excision

[Placeholder — to be written together: the excision method. A girth-safe BFS tree of depth $\lfloor \frac{g-1}{2} \rfloor$ is removed from a known (k, g) -graph, leaving boundary vertices of deficient degree; repairing them with girth-safe edges yields a smaller (k, g) -graph. Include worked examples in the style of the earlier chapters.]

7.3 Implementation

The two edits in this chapter share one labor split: a learned or heuristic search proposes moves and exact checks decide. Refinement is a tabu search that, each iteration, embeds the current graph once and scores every candidate 2-swap from that shared embedding, applies the best non-tabu swap that lowers the cost f , and escalates to a 3-swap only when no 2-swap helps. Excision instead removes a tree and repairs the degrees it broke: the removed tree leaves deficient vertices, and k -regularity is restored by adding edges among those vertices alone, where an edge (u, v) is girth-safe only when u and v lie at distance at least $g - 1$, so the cycle it closes is no shorter than g . The repair is a backtracking search that always extends the most-constrained deficient vertex first.

```

func refine(G: Graph, g: Integer, budget: Integer):
    tabu: Queue of Swaps = empty;
    best: Graph = G;
    repeat up to budget:
        // f(G) counts the short cycles; 0 means girth has reached g
        if f(G) is 0:
            return G;
        embed G once with the evaluator network;
        candidates: List of Swaps = all degree-preserving 2-swaps of G;
        score every candidate from its endpoint embeddings in one batch;
        sort candidates by predicted drop in f, largest first;
        // the learned score only ranks; the true f is checked before
    use
        chosen: Swap = first candidate not in tabu whose true f is below
    f(G);
        if chosen exists:
            apply chosen to G;
            tabu <- chosen;
            drop the oldest tabu entry if the list is too long;
        else:
            // no 2-swap helps: try the larger move, checked exactly
            chosen = best 3-swap that lowers f;
            if chosen exists:
                apply chosen to G;
            if f(G) < f(best):
                best = G;
    return best;

```

Algorithm 6: Tabu search for edge-swap refinement, with learned ranking of 2-swaps.

```

// deficit maps each deficient vertex to the number of edges it still
needs
func repair(H: Graph, deficit: Dictionary of Vertex to Integer, g:
Integer):
    if deficit is empty:
        return H;
    u: Vertex = deficient vertex with the fewest girth-safe partners;
    partners: List of Vertices =
        deficient vertices v != u with distance from u to v in H at
least g - 1;
    if partners is empty:
        return failure;
    for each partner v in partners, least-constrained first:
        add edge (u, v) to H;
        reduce deficit at u and at v;
        result = repair(H, deficit, g);
        if result is not failure:
            return result;
        // undo and try the next partner
        remove edge (u, v) from H;
        restore deficit at u and at v;
    return failure;

```

Algorithm 7: Girth-safe repair after excising a tree, by backtracking on the deficient vertices.

8 Forge

[Placeholder — to be written: the Forge pipeline, which chains the methods of the previous chapters. A voltage search (currently the reinforcement-learning variant) produces candidates and is deliberately allowed to pass near-perfect graphs whose girth is just below g ; edge-swap refinement (Chapter 7) then closes the remaining gap; and excision (Chapter 7) reduces the vertex count. Describe the stage hand-offs, the design rationale for the near-perfect hand-off, and why this ordering. Reference Chapters 6 and 7 rather than re-explaining the components.]

8.1 Implementation

The three stages are chained into a single loop. Whenever the voltage search yields a k -regular lift, we branch on its girth: a lift already at girth g goes straight to excision; a near-miss whose girth is just below g is handed to edge-swap refinement to close the gap, and only if refinement succeeds does it continue to excision; a lift further than a small margin below g is discarded, since refinement is not meant to bridge a large gap. Excision then shrinks a valid graph as far as it stays a valid (k, g) -graph. We keep the smallest valid graph found, and a surrounding harness repeats the loop in parallel to chase records.

```
func forge(g: Integer, margin: Integer):
    // smallest valid (k, g)-graph found so far
    best: Graph = none;
    repeat:
        candidate: Graph = next k-regular lift from the voltage search;
        gc: Integer = girth of candidate;
        // too far below the target: refinement is not meant to bridge
it
        if gc < g - margin:
            continue;
        // near-miss: close the small gap by refinement
        if gc < g:
            candidate = refine candidate toward girth g;
            if girth of candidate < g:
                continue;
        // candidate is now a valid (k, g)-graph; shrink it
        candidate = excise candidate while it stays a valid (k, g)-
graph;
        if best is none or candidate has fewer vertices than best:
            best = candidate;
    return best;
```

Algorithm 8: The Forge pipeline: voltage search, near-miss refinement, then excision.

9 Results

This chapter reports the final evaluation of every method developed in the previous chapters on a common footing: the per-node prediction probes (degree and minimum cycle) and the (k, g) -graph construction methods (classical search, direct reinforcement learning, voltage-lift search and its learned variants, and the forge cascade).

All neural models were trained on the PERUN supercomputer at the Technical University of Košice. The evaluation reported here, by contrast, was run entirely on CPU: every model is loaded and queried, and every construction search is executed, without a GPU. This reflects how the methods would be used in practice, where a single (k, g) target is attacked on commodity hardware, and it keeps the comparison between learned and purely combinatorial methods fair, since the latter never use a GPU either. For each architecture only the best trained model is evaluated.

The construction methods are compared on a fixed grid of 22 targets whose Moore bounds range from 6 to 106, spanning both the dense low-girth regime (where the cage is a small complete or complete-bipartite graph) and the sparse high-girth regime (where even existence is hard). Each stochastic method is run with 100 independent random seeds per target, and every attempt is given a wall-clock budget of 60 seconds. Graph order is reported relative to the Moore bound: a ratio of 1.0 means the method produced a graph of cage size, while a ratio of 2.0 means it produced a valid (k, g) -graph twice as large as the theoretical minimum.

9.1 Predicting Degree and Minimum Cycle

The two probing tasks isolate the two constraints that define a (k, g) -graph. Degree is a local counting property; the length of the shortest cycle through a vertex is a global one. Each trained architecture is evaluated on the same battery of 24 structured and random graphs (453 vertices in total), and scored by exact-rounded accuracy and mean absolute error (MAE).

Degree prediction is essentially solved by the right architecture. Table 1 shows that SAGE reaches perfect accuracy with under 5000 parameters, and the GPS–GIN hybrid is within rounding of it. The degree-normalised GCN variants are weakest, consistent with the expectation that dividing each message by neighbourhood size discards the very count the task asks for. The Loopy models are an instructive failure: they reach moderate accuracy but their regression head is numerically unstable, producing a MAE above 10^5 on the vertices it mispredicts.

Architecture	Parameters	Accuracy	MAE
SAGE	4 900	1.000	0.00
GPS–GIN	43 986	0.998	0.00
GPS–SAGE	43 854	0.980	0.02
GIN	151 948	0.887	0.34
Loopy	620 977	0.512	5.3×10^5
GPS–GCN	39 758	0.457	1.03
GCN	2 692	0.382	1.17

Table 1: Degree prediction: best model per architecture, sorted by accuracy.

Minimum-cycle prediction is much harder and is not solved by any architecture. Table 2 shows that the best model, the cycle-aware Loopy network, reaches only 0.510 accuracy, so it still mislabels almost half of all vertices. The ranking also reverses: SAGE and the GIN-based models, which dominate degree prediction, are among the weakest here, while the GCN family and Loopy — comparatively poor at degree — do relatively better on the cyclic task. This confirms that no single architecture is uniformly best, and that girth-related structure is not reliably recoverable as a supervised per-node label. It is this negative result that motivates treating girth through explicit search rather than direct prediction in the construction methods that follow.

Architecture	Accuracy	MAE
Loopy	0.510	1.45
GPS–GCN	0.311	1.43
GCN	0.227	1.80
SAGE	0.119	1.97
GPS–GIN	0.068	3.91
GPS–SAGE	0.068	3.97
GIN	0.055	16.18

Table 2: Minimum-cycle prediction: best model per architecture, sorted by accuracy.

Figure 6 summarises both tasks side by side. The contrast between the two bars for each architecture is the central observation: a model can master the local property while being almost useless on the cyclic one, and the only architecture that is even moderately balanced across the two is Loopy.




Figure 6: Per-node accuracy by architecture: degree vs. minimum cycle.

9.2 Constructing (k, g) -Graphs

The construction methods divide into three families. The classical searches (random walk, A^* , brute force) build a graph directly on a fixed vertex budget. The voltage-lift methods search for a voltage assignment whose lift is k -regular by construction and tune it for girth; this includes the purely algebraic tabu search, its two learned variants guided by the girth predictor and the tabu-cost predictor, and the reinforcement-learning lift policy (voltage_rl). The forge cascade chains a voltage producer with refinement and excision and is treated separately in the next section.

Table 3 reports, for each method, the fraction of the 2200 attempts that produced a valid (k, g) -graph, the mean order of the graphs found relative to the Moore bound, and the mean time per solved attempt.

Method	Success	Size / Moore	Time (s)
voltage_rl	0.799	1.98	1.26
voltage (girth predictor)	0.769	2.01	9.38
voltage (algebraic)	0.740	2.06	4.53
voltage (tabu predictor)	0.714	1.98	6.28
A*	0.545	1.01	1.40
random walk	0.500	1.00	0.65
forge (best producer)	0.462	1.27	3.97
brute force	0.136	2.00	1.06
direct RL	0.083	1.26	24.30

Table 3: Cage construction by method (22 targets, 100 seeds).

The dominant pattern is a trade-off between coverage and optimality, shown in Figure 7. The classical searches sit at a size ratio of essentially 1.0: when random walk or A* succeeds, it returns a graph of cage order. But they only succeed on the easier half of the grid and fail completely once the target tightens. The voltage family is the mirror image: it solves the most targets — voltage_rl succeeds on 80% of attempts and is also the fastest of the family — but the graphs it returns are about twice the Moore bound. A voltage lift trades graph order for the guarantee of regularity, so it reaches hard targets at the cost of size. The two learned voltage predictors do not change this picture materially: their success rates and graph sizes track the algebraic search closely, indicating that on these targets the learned guidance neither helps nor hurts the underlying tabu search appreciably.

Direct reinforcement learning is the weakest construction method by a wide margin (8% success, and the slowest per attempt). The policy loads and runs, but it does not learn to assemble cage-sized graphs, which is consistent with the difficulty the probing tasks already exposed for learning global structure directly.




Figure 7: Cage construction: coverage versus graph size relative to the Moore bound.

Three targets — $(4, 7)$, $(4, 8)$ and $(5, 7)$, with Moore bounds 53, 80 and 106 — were not solved by any method within the budget, and $(7, 5)$ was solved exactly once across all seeds. These are the largest targets on the grid, and they mark the practical reach of the methods as configured here.

The best Forge construction. Forge is the only method that actively minimises the graph it returns: it draws a k -regular lift from a voltage producer and then repeatedly excises and repairs to shrink it toward the cage. It exposes two independent choices — which producer drives the voltage stage, and whether the refinement and excision stages are enabled — and this section isolates both.

The producer choice matters most. Table 4 compares the four producers with the full pipeline enabled. The three non-reinforcement producers — the algebraic tabu search and its two learned variants — are statistically indistinguishable, all reaching about 46% success at a size ratio of 1.27. The reinforcement-learning producer is clearly worse on both axes: its success rate collapses to 31% and its graphs are larger. The same producer that was the single best standalone construction method (voltage_rl in Table 3) is the worst choice inside forge, because forge benefits from a diverse stream of near-cage lifts to excise, which the deterministic-leaning policy does not provide. On individual targets the gap is stark: with the tabu or girth producer forge solves $(4, 5)$ and $(4, 6)$ on essentially every seed, whereas the reinforcement producer succeeds on only about 1% of them.

Forge producer	Success	Size / Moore
voltage (tabu predictor)	0.462	1.27
voltage (girth predictor)	0.455	1.26
voltage (algebraic)	0.455	1.27
voltage_rl	0.310	1.52

Table 4: Forge by voltage producer, full pipeline, over all targets and seeds.

The excision and refinement stages were isolated by disabling each in turn. Excision is what makes forge a minimiser: with excision removed, the mean size ratio rises from 1.52 to 2.05, back to the level of an unshrunk voltage lift. Refinement, on the other hand, is counterproductive on these targets: removing it raises the success rate from 0.31 to 0.68, because the refiner spends the budget trying to close the girth gap of near-miss lifts that it usually cannot fix, instead of yielding to the cheaper excision path. The best forge configuration is therefore the full excision stage driven by an algebraic or tabu producer, with refinement playing at most a minor role.

Figure 8 shows the producer comparison on both axes. The payoff of this configuration is visible against Table 3: forge returns graphs at about 1.27 times the Moore bound, far closer to cage size than the ~ 2.0 of the standalone voltage methods, and on the small targets it reaches the cage exactly — it returns the 10-vertex $(3, 5)$ cage, the 14-vertex $(3, 6)$ Heawood graph, and meets the Moore bound on $(4, 6)$. Forge thus recovers much of the size optimality of the classical searches while retaining a large part of the voltage family’s reach, which is exactly the gap the other methods leave open.




Figure 8: Forge by producer: success rate and mean graph size vs. the Moore bound.

9.3 Summary

Across the two prediction probes, degree is solved exactly by a tiny SAGE model while the minimum-cycle length resists every architecture, which steers the construction methods

toward explicit girth search. Across the construction methods, no single approach is best on every axis: the voltage family reaches the most targets but oversizes its graphs, the classical searches are size optimal but narrow, and forge — run with the full excision stage and an algebraic or tabu producer — is the most balanced, returning near-cage graphs on a substantial portion of the grid. The largest targets on the grid remain out of reach for all methods within the budget tested.

10 Future Work

[To be written near the end. Candidate directions to cover: stronger voltage-lift search (pruning, canonicalization) before learning is applied, better training signals and hard-negative mining for the learned scorers, rewards that favor minimality (smaller lift orders / cage-finding) rather than just feasibility, training the excision-repair policy properly and adding excision as a full construction method, and learning as move-ranking inside existing search rather than end-to-end generation.]

11 Conclusion

[To be written last, once Chapter 9 is settled. Should summarize: the question (can GNNs help generate (k, g) -graphs), the preparatory finding (local vs. cycle-related properties), the construction methods tried (direct RL, voltage lifts and learned variants, edge-swap refinement), and the overall verdict on where learning helps relative to structured non-learned search. Keep the final claim aligned with what Chapter 9 actually shows.]

Bibliography

- [1] G. Exoo and R. Jajcay, “Dynamic Cage Survey,” *The Electronic Journal of Combinatorics*, 2013, doi: 10.37236/37.
- [2] G. Exoo, J. Goedgebeur, J. Jooker, L. Stubbe, and T. Van den Eede, “New small regular graphs of given girth: the cage problem and beyond.” 2025.
- [3] M. Grohe, “The Logic of Graph Neural Networks,” in *Proceedings of the 36th Annual ACM/IEEE Symposium on Logic in Computer Science (LICS)*, 2021, pp. 1–17. doi: 10.1109/LICS52264.2021.9470677.
- [4] V. K. Garg, S. Jegelka, and T. Jaakkola, “Generalization and Representational Limits of Graph Neural Networks,” in *Proceedings of the 37th International Conference on Machine Learning (ICML)*, 2020, pp. 3419–3430.
- [5] J. Gilmer, S. S. Schoenholz, P. F. Riley, O. Vinyals, and G. E. Dahl, “Neural Message Passing for Quantum Chemistry,” in *Proceedings of the 34th International Conference on Machine Learning*, 2017, pp. 1263–1272.
- [6] T. N. Kipf and M. Welling, “Semi-Supervised Classification with Graph Convolutional Networks,” in *International Conference on Learning Representations*, 2017.
- [7] W. L. Hamilton, R. Ying, and J. Leskovec, “Inductive Representation Learning on Large Graphs,” in *Advances in Neural Information Processing Systems*, 2017.
- [8] K. Xu, W. Hu, J. Leskovec, and S. Jegelka, “How Powerful are Graph Neural Networks?,” in *International Conference on Learning Representations*, 2019.
- [9] W. Hu *et al.*, “Strategies for Pre-training Graph Neural Networks,” in *International Conference on Learning Representations*, 2020.
- [10] L. Rampášek, M. Galkin, V. P. Dwivedi, A. T. Luu, G. Wolf, and D. Beaini, “Recipe for a General, Powerful, Scalable Graph Transformer,” in *Advances in Neural Information Processing Systems*, 2022.
- [11] R. Paolino, S. Maskey, P. Welke, and G. Kutyniok, “Weisfeiler and Leman Go Loopy: A New Hierarchy for Graph Representational Learning,” in *Advances in Neural Information Processing Systems*, 2024.
- [12] A. Lubotzky, R. Phillips, and P. Sarnak, “Ramanujan graphs,” *Combinatorica*, vol. 8, no. 3, pp. 261–277, 1988, doi: 10.1007/BF02126799.
- [13] G. Exoo and R. Jajcay, “On the girth of voltage graph lifts,” *European Journal of Combinatorics*, vol. 32, no. 4, pp. 554–562, 2011, doi: 10.1016/j.ejc.2010.12.003.

- [14] G. Araujo-Pardo, C. Balbuena, and J. C. Valenzuela, “Constructions of bi-regular cages,” *Discrete Mathematics*, vol. 309, no. 6, pp. 1409–1416, 2009, doi: 10.1016/j.disc.2008.02.012.
- [15] G. Exoo, R. Jajcay, and T. Raiman, “On decreasing the orders of (k, g) -graphs”, *Journal of Combinatorial Optimization*, vol. 46, p. 26, 2023, doi: 10.1007/s10878-023-01092-9.
- [16] D. Selsam, M. Lamm, B. Bünz, P. Liang, L. de Moura, and D. L. Dill, “Learning a SAT Solver from Single-Bit Supervision,” in *International Conference on Learning Representations*, 2019.
- [17] G. Bouritsas, F. Frasca, S. Zafeiriou, and M. M. Bronstein, “Improving Graph Neural Network Expressivity via Subgraph Isomorphism Counting,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 45, no. 1, pp. 657–668, 2023, doi: 10.1109/TPAMI.2022.3154319.
- [18] A. Z. Wagner, “Constructions in combinatorics via neural networks.” 2021.
- [19] I. Damjanović, U. Milivojević, I. Đorđević, and D. Stevanović, “RLGT: A reinforcement learning framework for extremal graph theory.” 2026.
- [20] B. C. A. Freire, N. Delfosse, and A. Leverrier, “Optimizing hypergraph product codes with random walks, simulated annealing and reinforcement learning,” in *2025 IEEE International Symposium on Information Theory (ISIT)*, 2025.
- [21] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal Policy Optimization Algorithms.” 2017.
- [22] A. Y. Ng, D. Harada, and S. Russell, “Policy invariance under reward transformations: theory and application to reward shaping,” in *Proceedings of the Sixteenth International Conference on Machine Learning (ICML)*, Morgan Kaufmann, 1999, pp. 278–287.